

Additional Benchmark Experiments

By Xiufeng Liu

Some Directions

- Data File Format
- Data Compression
- Data Serialization Methods
- Index
- File System
- Shark (on Spark) vs. Spark

Data File Format

- Data file formats:
 - Text file format
 - SequenceFile format (<http://blog.cloudera.com/blog/2011/09/snappy-and-hadoop/>)
 - Binary storage format can improve Hadoop scan performance by 3X [1] than text format
 - Parquet: columnar data format [2]
 - RCFile: record columnar file format [3]
 - Skip list column format [1]
- [1] Floratou, A., Patel, J. M., Shekita, E. J., & Tata, S. (2011). Column-oriented storage techniques for MapReduce. Proceedings of the VLDB Endowment, 4(7), 419-429.
- [2] Parquet. www.parquet.io
- [3] He, Yongqiang, et al. "RCFile: A fast and space-efficient data placement structure in MapReduce-based warehouse systems." Data Engineering (ICDE), 2011 IEEE 27th International Conference on. IEEE, 2011.

Data Compression

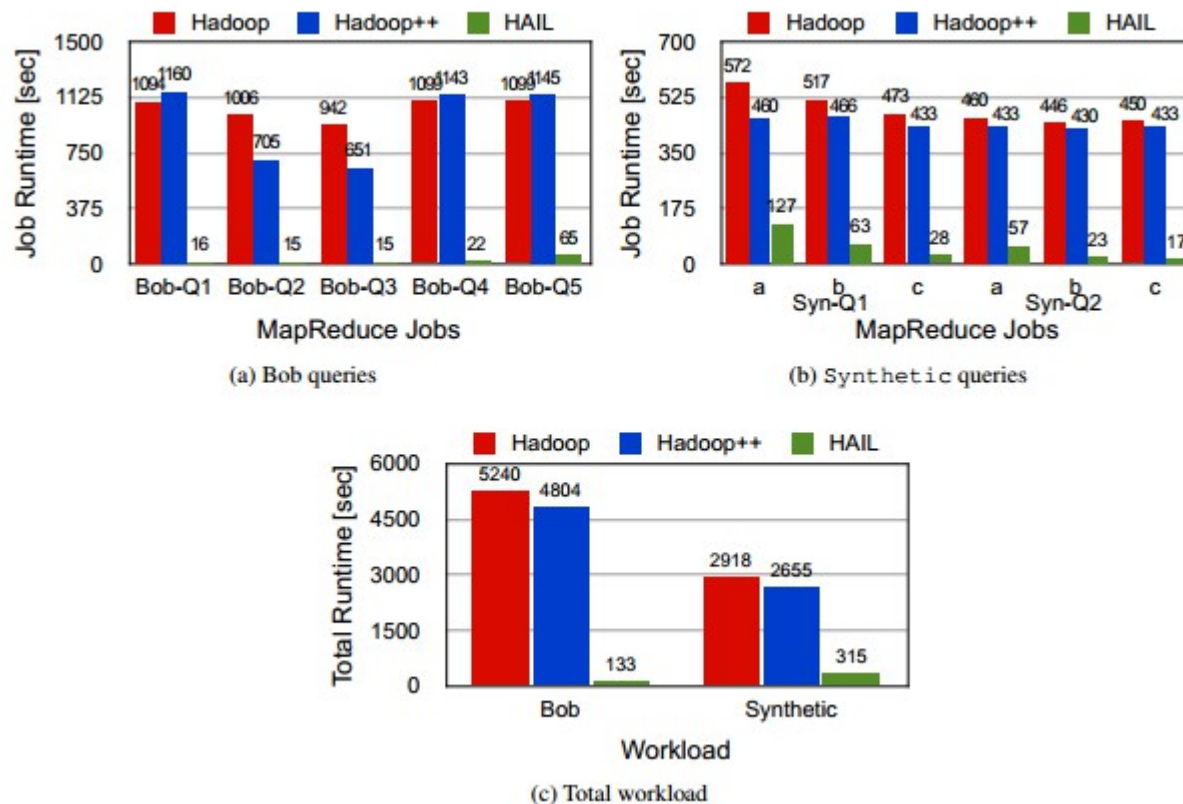
Format	Algorithm	Strategy	Emphasis	Comments
zlib	Uses DEFLATE (LZ77 and Huffman coding)	Dictionary-based, API	Compression ratio	Default codec
gzip	Wrapper around zlib	Dictionary-based, standard compression utility	Same as zlib, codec operates on and produces standard gzip files	For data interchange on and off Hadoop
bzip2	Burrows-Wheeler transform	Transform-based, block-oriented	Higher compression ratios than zlib	Common for Pig
LZO	Variant of LZ77	Dictionary-based, block-oriented, API	High compression speeds	Common for intermediate compression, HBase tables
LZ4	Simplified variant of LZ77	Fast scan, API	Very high compression speeds	Available in newer Hadoop distributions
Snappy	LZ77	Block-oriented, API	Very high compression speeds	Came out of Google, previously known as Zippy

Data (De-)Serialization

- Readable/Writable (built-in support by Hadoop)
- Avro (<http://avro.apache.org/>)
- Thrift (<https://thrift.apache.org/>)
- Protobufs (<https://code.google.com/p/protobuf/>)

Indexing on HDFS

- Hadoop++
- HAIL (Lazy Indexing and Adaptivity in Hadoop)



Reference:

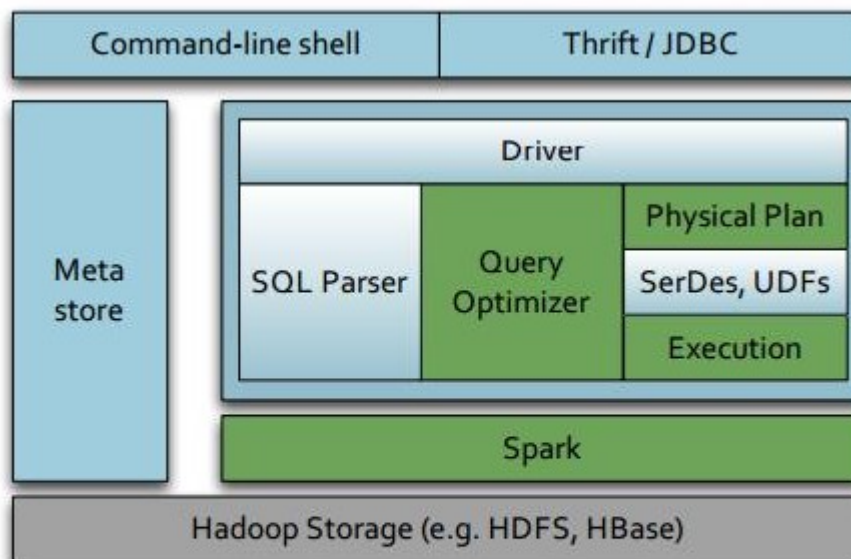
- Dittrich, Jens, et al. "Hadoop++: Making a yellow elephant run like a cheetah (without it even noticing)." Proceedings of the VLDB Endowment 3.1-2 (2010): 515-529.
- Dittrich, Jens, et al. "Efficient OR Hadoop: Why not both?" see <http://sites.computer.org/debull/A13mar/jens.pdf>

File System

- Hadoop Distributed File System (HDFS)
 - Data co-location vs. no co-location (see CoHadoop)
- Network File System
- Access file on disk/HDFS VS. cached files (e.g., ramfs, tmpfs)

Shark on Spark

- Shark builds on Hive that translate SQL-like queries to execution plans running on Spark
- Our Purpose: Present our experience of using Shark of facilitating processing time-series data vs. merely using Spark
- Benefits and comprises?



Shark Architecture